



Figure Z — LLM Cost and Utilisation Analysis

Real campaign data (Parts A) + Estimated tokens (Part B) + API pricing estimate (Part C) + Latency logged (Part D)

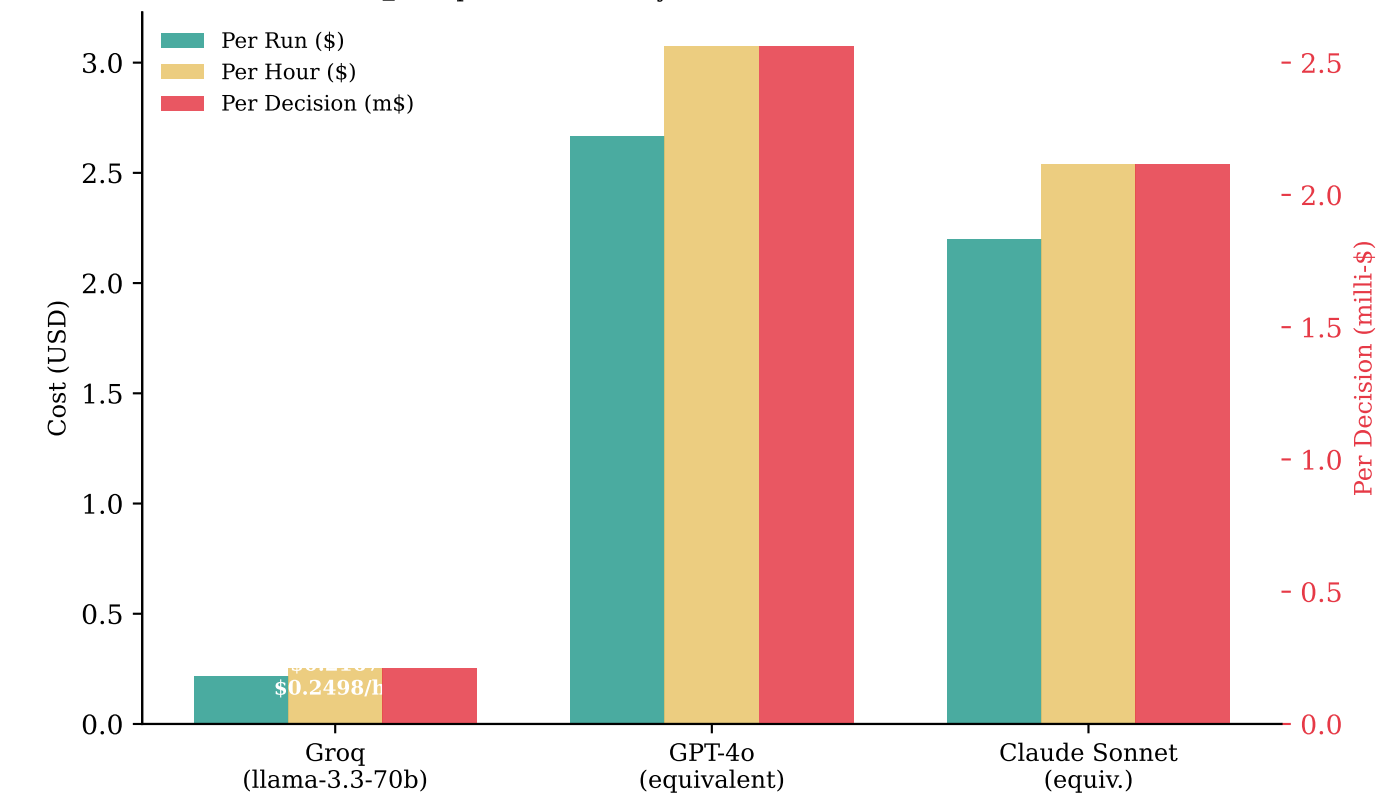
A — LLM Invocation Statistics
14 runs | LLM utilisation: 100% | Action rate: 44.6%



B — Token Usage Distribution
Mean=320 Median=319 p95=465 Max=480
[Source: dataset_agentic_medium.csv | ESTIMATED token counts]



C — API Cost Estimation
[Estimated | 320 tokens/decision | 1041 cycles | 70/30 in/out split]
□ Groq was used locally — no actual API cost incurred



D — Decision Latency Breakdown
Mean=232ms p95=370ms Max=385ms LLM share=91%
[Stage decomposition: architectural estimate | Total = real logged]

